

Old Divides, New Technology: A Video Experiment on AI Framing and Policy Preferences (Working Paper)

Julian Haddad

October, 2025

Abstract


Artificial Intelligence will transform labor markets by complementing some workers while displacing others. Using a preregistered video survey experiment ($N \approx 1,000$) of U.S. white-collar workers, I examine how messaging about AI as either a job threat or productivity tool shapes policy preferences on redistribution and regulation. I find voters are more responsive to threat frames than optimism frames: AI job threat messaging lowers expected personal income and increases pessimism about aggregate employment. Notably, Democrats prove more receptive to optimism frames about AI complementarity, while Republicans respond more strongly to substitution threats. However, both groups revert to traditional partisan policy preferences on redistribution and regulation. When faced with AI threat frames, those with low AI knowledge, low AI self-efficacy, less education, or downward social mobility prefer "predistribution" in the form of AI regulation over wealth redistribution. Despite AI's disruptive potential, voters' responses to AI threats follow predictable partisan and class patterns, with the most vulnerable often preferring regulation over redistribution.

Introduction

We are now in the midst of a new wave of technological change, that of Artificial Intelligence (AI). This technology is simultaneously accompanied by pessimistic messaging about AI replacing jobs *and* optimism about how it will empower workers. On the side of pessimism and in a rare moment of bipartisan agreement, both Bernie Sanders and Tucker Carlson have called for regulating AI-powered self-driving trucks, warning that widespread job displacement would devastate the economy.¹ Yet many of the workers potentially affected by AI also hear a starkly different message: that AI tools like ChatGPT will make them more productive, help them earn more, and enhance rather than threaten their careers.

Consider the following image from a [2025 Washington Post interactive](#) that invited readers to search their occupation and learn whether AI would likely "automate" or "augment" their job, a fitting illustration of how white-collar workers now face competing, simultaneous narratives about the same technology.

¹ [Bernie Sanders argues self-driving trucks need restrictions \(2025\).](#)
[Tucker Carlson calls to ban self-driving trucks \(2018\).](#)



Unlike previous automation waves that predominantly affected factory workers with limited dual-use potential (i.e., a laid off automobile factory worker was unlikely to buy their own CNC machine and get wealthy building cars using new automation tech), AI presents white collar workers with two possible realities: AI will easily replace many workers, but it can also help workers learn new, valuable skills or simply make their jobs easier to do.

How does framing AI as either a job threat or a productivity tool shape voters' policy preferences? The question matters because pessimistic economic expectations are known to strengthen preferences for redistribution and regulation (Benabou and Ok 2001; Thewissen and Rueda 2019), and because AI's potential to displace college-educated white-collar workers represents a departure from 20th-century automation shocks that primarily affected less-educated manufacturing workers or clerical workers. In fact, it threatens the exact workers who have traditionally been safe from technological replacement. If different workers rationally anticipate becoming "AI winners" or "AI losers," we might expect a bifurcation in policy preferences, with threatened workers demanding redistribution and regulation while beneficiaries oppose both.

Yet existing research offers conflicting predictions. Some studies find that perceived job risk from AI increases support for redistribution (Golin and Rauh 2024), while others find this effect only among those with high AI knowledge (Heinrich and Witko 2025). Importantly, no prior research has experimentally tested whether optimistic frames about AI complementarity can reduce support for redistribution, nor whether the effects of pessimistic substitution frames vary systematically by individual characteristics like AI self-efficacy, partisanship, or occupational exposure to AI risk.

To address these gaps, I fielded a [preregistered](#) video survey experiment ($N \approx 1,000$) among U.S. white-collar workers aged 25-54, randomly assigning respondents to view short videos framing AI as either substituting white-collar labor (job threat) or complementing it (productivity boost). I then measured preferences on redistribution (general support, taxing the rich, Universal Basic Income), regulation (banning AI layoffs, mandating safety measures, creating new oversight agencies), and labor protection (job retraining, union membership).

I find three main results. First, voters are more responsive to threat frames than optimism frames: the substitution treatment significantly reduces expected personal income and increases pessimism about aggregate employment, while the complementarity treatment produces no significant effects in the aggregate. Second, partisan asymmetries emerge in receptivity but not policy responses: Democrats prove more receptive to optimism frames about income gains, while Republicans respond more strongly to threat frames. Yet both groups revert to traditional partisan policy preferences, with Democrats supporting redistribution and Republicans opposing social safety net expansion regardless of treatment. Third, individual characteristics condition treatment effects in predictable ways: when faced with AI threat frames, those with low AI self-efficacy, low AI knowledge, no degree, or downward social mobility prefer "predistribution" through AI regulation over wealth redistribution, while those in high-risk occupations (measured via EPOCH exposure scores from Loaiza & Rigobon, [2024](#)) increase support for redistribution.

These findings carry important implications. First, that voters have a persistent AI pessimism which will likely color the policy responses they prefer. Second, that the trend of partisanship being an extremely powerful force in shaping voters' policy preferences will apply to AI preferences regardless of how AI is framed in the media. Indeed, voters are receptive to the messaging of AI as a looming economic problem but perhaps have party-informed philosophical differences about the appropriate role the state should play in solving that problem. Additionally, this study shows that regulation (a form of predistribution) may be a more politically salient and feasible policy response than redistribution, and that it is particularly preferred by those situated lower on the socioeconomic hierarchy (Kuziemko et al., [2023](#)).

The paper proceeds as follows. The next section reviews relevant literature on automation, political attitudes, and the emerging research on AI and policy preferences, highlighting theoretical gaps this study addresses. I then develop hypotheses about how AI framing should shape preferences, conditioned by individual characteristics. After describing the experimental design and data, I present results and conclude with implications for the political economy of technological change and the future of public policy in the age of AI.

Literature Review

Technological advancements have had profound impacts on politics throughout history, from the Cotton Gin precipitating economic conditions that led to the Civil War in the 19th century to the role of the assembly line in strengthening union politics in the early 20th century (Acemoglu and Johnson 2023; [Kerrissey & Schofer, 2018](#); [Streeck and Hassel, 2003](#)). Factory automation contributed to late 20th century deindustrialization in the West, (Acemoglu and Johnson 2023 pp. 385; [Alder et. al, 2014](#)), and the political reverberations of deindustrialization are difficult to overstate. Iversen contends that workers were so profoundly dislocated by the dramatic drawdown in employment in agriculture and manufacturing and growth of the service economy that they supported politicians who dramatically expanded the welfare state (Iversen, [2001](#), p.3). More recently, deindustrialization has likely contributed to the recent wave of populist election victories in the West beginning in the 2010s ([Alder et. al, 2014](#), [Broz et al., 2021](#), Colantone & Stanig 2018).

What relevant historical precedents exist with regard to AI? In the 20th century, automation predominantly posed a threat to routine jobs typically performed by less skill-intensive occupations with limited formal education (Autor et al., [2003](#), Spitz-Oener, [2006](#)). However, recent advancements in Generative Artificial Intelligence (gen-AI) could extend this threat to non-routine jobs, which are often filled by highly skilled, educated workers who are more politically active and thus may have a greater impact on policy (Berinsky and Lenz [2011](#), Hillygus [2005](#), Jacobs, [2024](#), Kim [2023](#)).

The nascent literature on generative AI's impacts on political preferences shows an unclear relationship between the threats of technological change to individuals' work and their political preferences. Using survey experiments with varying information treatments Golin and Rauh ([2022](#)) show that perceived job risk from AI increases support for redistribution, while Heinrich and Witko ([2025](#)) only find this to be the case among those who are highly knowledgeable about AI. Heinrich and Witko ([2025](#)) offer another important finding: that AI regulation is a key desire of those with high AI knowledge. This particular finding represents a slight departure from Kuziemko et al. 's ([2023](#)) finding that the less educated in the US electorate tend to support regulation while the more educated prefer redistribution.

Additionally, the results of the two aforementioned survey experiments are particularly puzzling as previous research on automation would have predicted that even workers who do not know much about how AI works should still respond to information treatments stating that AI will take their jobs with demands for welfare state expansion as insurance (Iversen 2001, Thewissen and Rueda 2017).

I believe that a potential explanation for these mixed findings is that most existing research does not take into account the expectation among some white-collar workers that AI will boost or complement their incomes. While a factory worker facing pressure from automation might not have seen new, capital intensive car manufacturing technologies as something that could potentially empower them to make more money, white collar workers face a different outlook. On one hand, they may see AI as a potential cause of future unemployment. On the other hand, obtaining access to AI technologies is relatively cheap (and often free), and can be used by white collar workers to boost their productivity and potentially to earn more money.

Corroborating this observation, [Green et al., \(2025\)](#) fielded a survey² in the United Kingdom that highlights the importance of the perception of AI being a benefit or a detriment to one's career. They find that many British workers anticipate either being an AI 'winner' (10%) or 'loser' (24%), defined as whether or not survey respondents said they expect AI to either increase or decrease their job prospects.

Interestingly, their findings contradict previous empirical findings that indicate increased automation risk is associated with increased support for redistribution (Busemeyer and Alexander, 2021; Thewissen and Rueda 2017) by finding that AI winners and losers have no differences in their support for redistribution when surveyed. Their findings also seem to contradict what Golin and Rauh (2022) show, that when faced with the knowledge that experts think one's job will be automated by AI, support for redistribution *does* increase. Lastly, their work does not speak to Heinrich and Witko's (2054) finding that AI knowledge is an important predictor of support for redistribution and regulation of AI.

Green et al., (2025) second set of findings that is relevant to this project is that AI winners and losers view active labor market policies differently, where AI winners in fact are more supportive than AI losers of government funded retraining of jobs while AI losers are more hostile to immigration. It may be the case that even the *perception* of a looming economic shock from AI will increase hostility to immigration.

² Note: Green et al., (2025) do not field a survey *experiment*. The authors do not ask respondents to quantify the financial impact of AI on their career.

Green et al., (2025) usefully demonstrate that perceptions about being an AI winner or loser are politically meaningful, but answers to a number of important questions remain unanswered. For one, it is unclear the extent to which these effects vary based on how impactful someone expects AI to be on their career. Does someone expect AI to double their income or increase it by 10%? Does someone expect AI to make them jobless or simply keep their wages stagnant? As indicated by a great deal of research in political economy, the extent to which someone expects to either lose money or gain money affects their political and economic attitudes (Benabou and Ok, 2001). Additionally, the mixed findings of Golin and Rauh (2022), Heinrich and Witko (2024), and Green et al., (2025) beg the question of how exactly AI knowledge, expected future income, and perception of being an AI winner or loser affect political preferences.

Lastly, Green et al., (2025) performed this survey in the United Kingdom, which has a more universal and generous welfare state than the United States. A study in the United States would offer additional utility to scholars as we have reasons to believe that living in a country with a larger welfare state ameliorates attitudes toward job threats (Bussemeyer and Alexander, 2021).

Why New Research is Needed

The existing literature reveals three important gaps. First, while prior research has examined how *threat* frames shape policy preferences (Golin and Rauh 2022; Heinrich and Witko 2024), no study has experimentally compared the effects of *optimistic* versus *pessimistic* framing about AI. This is consequential because white-collar workers simultaneously receive both messages, yet we do not know whether positive frames can offset threat perceptions or whether negativity bias dominates.

Second, findings about who responds to AI information remain contradictory. Heinrich and Witko (2025) find that only high-knowledge respondents increase support for redistribution and regulation when exposed to job threat information, while Golin and Rauh (2022) find broader effects. Green et al. (2025) find no difference in redistributive preferences between self-identified AI "winners" and "losers," contradicting both prior studies. These inconsistencies suggest that individual-level moderators such as AI self-efficacy, occupational exposure, or partisan identity may condition how workers translate economic information into policy demands.

Third, the literature has not adequately theorized or tested whether AI's dual-use nature creates distinct policy responses compared to previous automation shocks. If AI can both threaten and complement the *same* workers and those workers are responsive to AI messaging, then standard models predicting that objective job threat leads to welfare state expansion (Iversen 2005; Golin and Rauh 2022; Thewissen and Rueda 2019) may be incomplete in the AI era.

I address these gaps by experimentally manipulating exposure to substitution versus complementarity frames and testing how individual characteristics (AI knowledge, AI self-efficacy, occupational risk, social mobility, and partisanship) moderate treatment effects on multiple relevant policy outcomes (redistribution, regulation, labor protections). This design offers us a clearer picture of the political psychology of policy preferences among workers impacted by AI and has important implications about whether AI will bifurcate the electorate along new economic lines or reinforce existing partisan and class divisions.

Theory

Because policy attitudes on economic redistribution are often driven by ones' expected future income (EFI) (Bénabou and Ok [2001a](#), Checchi and Filippin [2003](#), Ravallion and Lokshin [2000](#)). I expect workers who view AI as a threat to their EFI to support more redistributive and regulatory policies, and those who view AI as a potential boon to their incomes to want less redistribution and regulation on AI. In line with this thinking, I also expect concerns about unemployment to drive attitudes toward redistribution and regulation of AI. Additionally, I expect that these economic policy preferences are moderated by a number of individual characteristics, such as individual occupational risk, risk aversion, and partisanship (Alesina & Giuliano [2009](#), Busemeyer & Alexander [2021](#), Buser et al. [2020](#), Wallsten and Park [2015](#)).

Main Treatment Effects

H1: Exposure to the Substitution treatment increases support for redistributive and protective policies relative to Control.

H2: Exposure to the Complementarity treatment decreases support for redistributive and protective policies relative to Control.

Moderating Factors

AI Knowledge and AI Self-Efficacy

Knowledge of AI has also been shown to be a key variable in predicting how political attitudes are shaped by AI. Awareness of generative AI's labor-substituting potential among those who are knowledgeable about AI should also prompt increased support for regulation (Heinrich & Witko [2024](#)) as it is likely already viewed by a significant portion of respondents as either a hazard or risk. Golin and Rauh ([2022](#)), for example, show that workers reported a 32%-50% chance of

losing their job to AI or a robot in the next five years, and that treatments have the potential to increase automation concern.³

H3: AI knowledge will moderate treatment effects: more knowledgeable participants will show larger shifts in preferences in the expected direction.

Many respondents now possess surface-level information about AI without necessarily feeling confident using it. I therefore include AI self-efficacy alongside AI knowledge to capture variation in perceived competence and comfort with AI tools, which is an arguably more behaviorally relevant measure of how voters translate information about AI into policy preferences (Wang and Chuang, [2024](#)). I therefore pre-registered the following hypothesis.

H4: AI Self-Efficacy will moderate treatment effects in the expected direction, with more risk averse and less AI self-efficacious individuals supporting more redistribution and regulation. (Lown 2011; Wang & Chuang, 2023)

Risk Aversion

Those who are risk averse generally support more redistribution, and those who are more comfortable taking risks are less supportive (Dohmen et al., [2011](#)). These findings are fairly intuitive: those who are comfortable taking risks may be more entrepreneurial and see themselves as having much to gain from an environment of lower taxes and less redistribution. I pre-registered the following hypothesis:

H5: Risk aversion will moderate treatment effects in the expected direction, with more risk averse individuals preferring more redistribution and regulation

Social Mobility

We expect that individuals' perceived trajectories of social mobility will influence how they interpret economic disruptions from AI. Research shows that expected future income and perceptions of personal mobility shape support for redistribution and the welfare state (Bénabou & Ok 2001; Thewissen & Rueda 2017). Those who perceive themselves as downwardly mobile may view AI as compounding existing insecurity and thus favor stronger protective policies, regardless of whether AI is framed as a threat or opportunity. Conversely, those on upward mobility trajectories may view AI optimistically and oppose redistributive interventions that could dampen future gains.

³ To be clear, the authors did not ask about generative artificial intelligence because it was not yet publicly available, so I theorize that individuals will feel similarly about generative artificial intelligence to the way that the author is responding she did about robots and AI in the article in question.

H6: Downward social mobility will increase support for redistributive and regulatory policies regardless of treatment, while upward mobility will dampen such support.

Partisanship

In addition to economic self-interest and individual labor market risk, voters' responses to AI may also be filtered through their partisan commitments. A large body of research finds that individuals interpret economic information through partisan lenses, particularly on redistribution and regulation (Bartels, [2006](#); Iyengar & Krupenkin, [2018](#)). Republicans and Democrats differ not only in their policy preferences but in their interpretations of economic threats, with Republicans generally more skeptical of state intervention and more optimistic about market-based solutions (Gelman et al., [2010](#)). I therefore expect that partisan identity will shape both the interpretation of AI-related framing and the policy responses it elicits. Specifically, Republicans may be more receptive to the optimism of the complementarity frame, while Democrats may be more attuned to substitution threats, given their longstanding policy alignments.

H7: Partisanship will moderate treatment effects, with Democrats showing stronger support for regulation and redistribution under substitution framing, and Republicans showing stronger support for deregulation and lower taxation under complementarity framing.

Occupational Substitution Risk and Augmentation Potential

Another important moderator of political responses to AI is occupational exposure to automation risk **and complementarity potential**. Longstanding political economy theories of technological change suggest that occupational location within the labor market shapes both material exposure to disruption and the interpretive lens through which new technologies are evaluated (Autor, [2003](#); Thewissen & Rueda, [2019](#)). However, no previous research has specifically explored both how measures of occupational risk and potential augmentation or complementarity interact with policy preferences.

I hypothesize that the effects of AI framing treatments will be moderated by a respondent's occupational risk level as determined by this linkage. Specifically, individuals in high-risk occupations should be more sensitive to the Substitution treatment and more likely to support redistributive and protective policies such as regulation and job retraining. In contrast, individuals in high-augmentation occupations may be more persuaded by the Complementarity frame and more resistant to redistributive interventions.

H8: Occupational risk from AI will moderate treatment effects.

H8a: Among those assigned to the Substitution treatment, individuals in high-risk occupations will show greater increases in support for redistribution and regulation than those in low-risk occupations.

H8b: Among those assigned to the Complementarity treatment, individuals in high augmentation occupations will show greater declines in support for redistribution and regulation than those in high-risk occupations.

Survey Design

Between August 25 and September 7, 2025, 1006 participants were recruited through Prolific, a widely used UK-based web-survey platform. Respondents then completed an IRB-approved, short Qualtrics survey with a median response time of 8 minutes and 31 seconds. Eligibility was restricted to white-collar workers in the U.S. aged 25–54. Screening was conducted through Prolific, and checked for accuracy via self-reports of occupation. The final sample of respondents closely reflects the demographic composition of the U.S. white-collar labor force.

To calculate the partisanship, gender, age, and demographic breakdown of the U.S. white-collar workforce, public data from the American National Election Survey and General Social Survey (GSS) were used.⁴ 41.5% of respondents identified as Democrats, 18.3% independents, 4.2% Other, and 36% as Republicans. Respondent ages ranged from 25 to 54, with the average age of a respondent being 39.5. The average income of respondents was \$4157/month, or approximately \$50,000 annually. 68% of respondents had a Bachelor's Degree or higher.

To enhance external validity, the survey employs short video treatments that simulate how individuals increasingly consume political and economic information online instead of text vignettes ([Mutz 2021](#)).

After informed consent, participants were randomly assigned to one of three conditions:

- 1) **Group 1: Control Group**. Randomly shown one of four informational videos that are not related to Artificial Intelligence. These videos were a [National Geographic video on giraffes](#), a [video about how to juggle](#), a [clip from a documentary about the Giza pyramids](#), and a [documentary clip about the Grand Canyon](#).

⁴ I constructed demographic tables for the US white-collar workforce using publicly available data from the BLS and ANES. For more information on this, see appendix table A.2.

- 2) **Group 2: Complementarity**: Shown a video highlighting how AI can save them time at work and make them more money. The video can be viewed at [this link](#), and was preceded by the following blurb highlighting AI-white-collar complementarity:

“Several studies published in 2023–2025 report that workers who adopt AI tools can complete tasks faster and take on higher-value work. These analyses estimate potential pay increases of 20–30% for white-collar workers, alongside growth in overall employment in AI-exposed occupations.

Please watch this short video showing how Artificial Intelligence (AI) tools like ChatGPT are helping white collar workers earn more money by helping them complete their tasks faster.

Once you finish watching the video, you will be able to go to the next page.”

- 3) **Group 3: Substitution**: This group watched a video showing how AI was 3) already leading to layoffs and could do a great deal of white collar work. The video can be viewed at [this link](#), and was preceded by this blurb highlighting AI-white-collar substitution:

“Several studies from 2023–2025 warn that companies may use AI to automate a substantial share of routine professional work. These analyses estimate that about 30% of white-collar workers could face wage reductions or job loss if firms replace tasks previously done by people with AI systems.

Please watch this short informational video showing how Artificial Intelligence (AI) platforms like ChatGPT, Claude, and Gemini are already causing layoffs and reducing wages for white collar workers because of AI's ability to do work cheaply and quickly.

Once you finish watching the video, you will be able to go to the next page.”

All videos were embedded in the Qualtrics platform and could not be skipped. A post-video attention check confirmed comprehension, and participants who failed either of two attention checks were excluded. Of the 1006 respondents, 12 failed the attention checks.

Dependent Variables

Immediately following the video, respondents completed a battery of outcome measures. The primary dependent variables were:

- 1) **Redistribution index** (an average of support for general redistribution, taxing the rich, and taxing corporations)⁵
- 2) **Regulation Index** (averaged across four Likert-style AI regulation items adapted from Heinrich and Witko, 2024)
- 3) **Job retraining** (support for funding job retraining for those who lose their jobs to AI)
- 4) **Interest in joining a union**
- 5) **Universal Basic Income**⁶

Universal Basic Income is held analytically separate from other measures of redistribution because of its current status as a relatively fringe policy proposal. Support for UBI did not correlate with support for other measures of redistribution, however, taxing the rich, taxing corporations, and support for redistribution are strongly correlated (Pearson's $R \sim .65$)

Independent Variables

To examine heterogeneity in treatment effects, the survey includes several pre-treatment measures in addition to basic demographics:

- 1) Partisanship
- 2) AI knowledge (measured via factual questions about AI, adapted from Heinrich and Witko, [2024](#))

⁵ The redistribution question read: "Some people feel that the government should increase taxes a lot and spend much more on health and social services such as unemployment benefits and low-cost healthcare programs like Medicaid. Other people feel that the government should cut taxes a lot and spend much less on health and social services. And other people have views somewhere in-between. Please use the slider to indicate where your views fall on this scale." (0-100 scale)

⁶ The UBI question read: "There is a discussion about the introduction of universal basic income. Universal basic income is a concept where everybody, regardless of whether they work or not, receives the same amount of money every month from the government. If universal basic income were to be introduced, how much should everybody receive per month? Put 0 If you don't want UBI to be introduced."

- 3) AI self-efficacy (a six-item scale adapted from Wang and Chuang, [2024](#))
- 4) Risk aversion (0–10 scale from Dohmen et al., [2011](#))
- 5) Social Mobility (captured by asking both one’s childhood economic class and one’s expected future class. If a person’s reported childhood class was “Upper Middle: but their expected future class is “Lower Class”, they’d be considered highly downwardly mobile)

In order to operationalize the potential for a respondent’s job to either be augmented by AI or be substituted by AI, I link respondents’ self-reported occupations (coded using Standard Occupational Classification (SOC) 2018 codes) to occupational automation exposure scores from the EPOCH dataset (Loaiza and Rigobon [2024](#)). The EPOCH framework estimates the substitutability of over hundreds of U.S. occupations based on their task content and the capabilities of emerging AI systems, particularly large language models and other generative AI technologies. This enables a granular, occupation-level measure of whether AI is likely to complement or substitute for a respondent’s job tasks, making it well-suited for testing how perceived versus objective exposure shapes political responses. Loaiza and Rigobon ([2024](#)) thus construct two variables of interest for this study:

- 1) Occupational AI Risk (a continuous measure between 0 and 1 where a higher value means the job is at greater risk of being replaced by AI)
- 2) Occupational AI Augmentation Potential (a continuous measure between 0 and 1 where a higher value means the job can be made even more productive by AI)

Additional post-treatment manipulation checks and mediator variables captured beliefs about personal and sociotropic economic risks from AI:

- 1) Expected income change due to AI
- 2) Perceived AI-related unemployment risk in the next five years
- 3) Expected impact of AI on total number of jobs in the economy

Results are quantified via t-tests and outcomes are normalized relative to the control group mean wherever possible for ease of interpreting effects sizes in terms of standard deviations from the control group mean. I seek treatment effect sizes of at least Cohen’s $d = .2$, the threshold for a small effect (Paul, 2010). An effect size of .35 would put the power of this study at a similar level to that of Heinrich and Witko (2025), who find effect sizes of 0.34 standard deviations on support

for regulating AI when those with high AI knowledge were treated with transformation of work treatments.

Main Findings

Consistent with prior studies, I find mixed support for the hypothesis that AI framing affects public policy preferences. In the aggregate, the substitution treatment produced directionally consistent but mostly non-significant increases in support for redistributive and regulatory policies, relative to the control group. I fail to reject the null for H1 and H2. The complementarity treatment had no statistically significant aggregate effects on any outcome, including redistribution, regulation, UBI, job retraining, or union interest.

Figure 1 shows treatment effects in terms of standard deviations from the control group mean. For ease of interpretation, outcomes are standardized relative to the control group (subtracting control mean, dividing by control SD). Coefficients therefore represent treatment effects in control-group standard deviation units. While most effects did not reach significance, the substitution treatment is directionally associated with greater support for redistributive, regulatory, and labor protectionist policies, with the exception of support for Universal Basic Income. This pattern suggests that AI threat frames may activate economic concerns, but these concerns do not consistently translate into broad-based public policy support in the full sample.

Mediation via Economic Expectations

I next tested whether the substitution and complementarity treatments influenced my three pre-registered mediators: expected future income, perceived unemployment risk, and beliefs about the effect of AI on the total number of jobs. The substitution treatment significantly reduced both personal income expectations (by 7.1%) and beliefs about job creation (sociotropic concerns), with small-to-medium effect sizes (Cohen's $d \approx -0.21^7$ and -0.40 , respectively).⁸ These findings support the theoretical claim that threat frames activate economic pessimism. The complementarity treatment had no significant effect on any mediator.

⁷ EFI is displayed in table 1 here as a percentage instead of a z-scored outcome for ease of interpretation.

⁸ See appendix section A.3 for full effect size reports.

Moderation by Individual Characteristics

Partisanship

Partisanship emerges as a strong moderator of framing effects, but in surprising ways. I pre-registered a hypothesis stating that Republicans would be more responsive to the Complementarity prime and less responsive to the Substitution prime with regard to our outcomes of interest, while the inverse would be true for Democrats. In fact, I find the *opposite to be true*. Democrats are in fact more receptive to cues highlighting how AI will increase their income than Republicans, despite having similar levels of education and monthly incomes.⁹




Figure 1

Surprisingly, Republicans were more responsive to the substitution treatment than Democrats in terms of shifting income expectations. Republican respondents in the control group expected a net income increase from AI of +13.8%, but this dropped to just +2.4% after the substitution treatment, representing a decline of more than 11 percentage points. Democrats, in contrast, showed a statistically significant 9-point increase in income expectations in the complementarity treatment.

⁹ Mean education levels are nearly identical across parties, with most respondents holding a BA, followed by graduate degrees. The mean monthly income of Republicans is ~\$4500 compared to ~\$4100 among Democrats




Figure 2

Given the increased AI pessimism observed among Republicans, one might expect increased support for policies like job retraining, general redistribution, or even UBI. Yet, the opposite is true. Among Republicans, I find the substitution treatment is associated with *less* support for UBI, and *less* support for job retraining. Democrats exposed to the substitution treatment, however, prefer more redistribution at borderline significance levels ($p = .056$).

These results may indicate some sort of reflexive falling back on age-old, party-specific economic prescriptions in the presence of information about the economy. When presented with

information about how their sector of the economy is changing, Democrats will choose the party prescription of increasing redistribution, while Republicans support reductions in the social safety net.

AI Knowledge

Drawing on Heinrich and Witko (2025), H3 states that AI knowledge strengthens the substitution effect on all dependent variables and weakens the effect of the complementarity treatment on all dependent variables. Instead, I find evidence for a different framework at play: that of those with low AI knowledge preferring “pre-distribution” in the form of regulation, while null effects are found for those with high AI knowledge.¹⁰




Figure 4

¹⁰ The AI Knowledge Quiz included the following four multiple choice questions:

- 1) Who is Sam Altman?
- 2) What does RLHF stand for?
- 3) What is “OpenAI”?, and
- 4) Large language models like ChatGPT and Claude are primarily trained to do what?

I dichotomously operationalize high AI knowledge as getting 4/4 AI quiz# questions correct, while low AI knowledge is 2 or less questions correct. The mean quiz score was 2.74/4. No other statistically significant effects are found with respect to knowledge. Mediation effects of AI knowledge are available in the appendix.

AI Self-Efficacy

In a pre-registered hypothesis, I proposed AI self efficacy would dampen demands for compensatory policies. I find evidence for those with low AI self-efficacy supporting greater AI regulation when primed with the AI job threat. Rather than seeking redistribution, these individuals prefer regulating AI: a form of what some Kuziemko et al (2023) call “predistribution.”




Figure 5




Figure 6

I measure AI Self Efficacy (AISE) using a shortened¹¹ version of the AISE instrument developed by Wang & Chuang (2024). When split into two subgroups, those above the median and those below the median level of AI self efficacy, I find that those with high AI self efficacy are more likely to support increased regulation of AI when primed with substitution, while those with low AI self efficacy are *less* likely to support job retraining in the substitution treatment.

In OLS models, I observe significant interaction effects: high self-efficacy respondents in the substitution condition showed decreased interest in union membership, while in the complementarity condition, high self-efficacy predicted reduced support for job retraining. These

¹¹ The original scale comes from a 22-item instrument. I shorten it to the following six questions, asked via a Likert scale.

- 1) When interacting with artificial intelligence (AI) technologies/products, I feel very Relaxed
- 2) I can happily interact with AI technologies/products smoothly.
- 3) AI technologies/products are good aids to learning.
- 4) AI technologies/products jargon does not baffle me.
- 5) AI technologies/products help me to save a lot of time.
- 6) I find it easy to get AI technologies/products to do what I want it to do.

"Not at all true" = 1,
 "Hardly true" = 2,
 "Moderately true" = 3,
 "Exactly true" = 4,
 "Not sure" = NA

The AISES index is the mean value of a respondents' answers to these questions.

effects underscore that self-efficacy more accurately conditions how individuals process AI-related messaging.

Education




Figure 7

In keeping with my findings with regard to AI self-efficacy and AI-knowledge, pre-distribution in the form of AI regulation is the preferred policy response of those with no degree in the substitution cohort, while those with some degree prefer UBI not be introduced.

Occupational Exposure

I pre-registered a prediction that occupational risk from AI would be associated with greater support for redistribution and regulation as a form of insurance against job threats, and that complementarity or augmentation would be associated with the inverse. I find support for the notion that treatments are able to cue respondents to “rationally”¹² consider their own personal financial prospects with regard to redistribution, but not regulation. These results suggest that the complementarity treatment may be able to trigger notions of upward mobility and pocketbook concerns that, in turn, lead respondents to support lower taxes on the wealthy given that they expect it may benefit them.

To operationalize AI risk and AI augmentation, I employ the *risk* and *augmentation* metrics from the EPOCH dataset created by Loaiza and Rigobon (2024), which ranges from 0 to 1, where a higher score means higher risk or higher augmentation potential. The average AI risk score for our sample is .55. I do not observe systematic variation by party, as Republicans and Democrats have nearly identical average risk and augmentation scores, with the average risk scores for

¹² In the narrow sense of a rational economic actor who is seeking to maximize their financial benefits.

Republicans and Democrats being $\sim.56$ and $\sim.54$, respectively, and both average augmentation scores being $\sim.4$.¹³

Significant Effects Only — Complementarity $\times$ Risk




Figure 8

Social Mobility

I pre-registered an exploratory hypothesis that downward social mobility would increase support for our outcomes of interest as a form of income protection. I am able to reject the null hypothesis with regard to regulation preferences, but not with regard to redistribution, tax policy, or union interest.

As much of this study's theory relies on how individuals' upward or downward economic prospects informs their preferences on economic policies, I controlled for upward or downward social mobility by asking respondents what their childhood socioeconomic status was (lower class, lower middle class, middle class, upper middle class, upper class) and what they anticipated their future status to be. If they made a jump of at least two classes, they are considered highly upwardly mobile, while those who made a jump *downward* of at least two classes are considered highly downwardly mobile.

¹³ 94.6% of respondents were matched to an EPOCH risk or augmentation score using SOC18 codes. The remaining $\sim 5\%$ selected professions for which there was no matching EPOCH score.




Figure 9

Among the highly upwardly mobile, I do not observe statistically significant results at the $p < .05$ level, but those who were highly downwardly mobile support greater regulation of AI, regardless of the treatment group they were in. I do not observe any changes in other policy outcomes of interest. These results imply that among some, the framing of AI as a helper or a substitute to their labor won't matter in forming some policy preferences, as they already consider their economic future to be gloomy.

The complementarity treatment, however, produces an expected reaction on those of stable social mobility by reducing their support for taxing corporations ($p < .05$) and, with borderline significance ($p < .10$) reducing support for redistribution, potentially indicating that optimism about AI's benefits may only be salient among those with secure trajectories.

These results support H6 and bolster the claim that AI is interpreted through the lens of one's *perceived* trajectory, not just objective exposure.

Discussion

AI is transforming the workplace for white-collar workers, many of whom are exposed to both optimistic and pessimistic messaging. Despite widespread claims that AI can boost productivity and wages, this study finds voters are more responsive to job threat framing than to promises of AI augmentation, suggesting a persistent AI pessimism. Threat cues significantly reduce expected income and increase pessimism about the labor market, but these shifts in belief do not reliably translate into support for redistributive policy.

Instead, the translation of economic beliefs into policy preferences appears to follow existing partisan alignments and a preference for predistribution by the less knowledgeable and less educated over redistribution. Voters with low AI self-efficacy or downward mobility are more supportive of regulating. Despite the differences in how each party expects AI to affect their pocketbooks, partisans largely support typical economic prescriptions of their parties: less support on social safety nets for Republicans, and increased redistribution for Democrats.

If each party represents fundamentally different philosophies about the role the state should play in the economy and how it should respond to economic challenges, it stands to reason that Democrats and Republicans will support their party's standard economic policy prescriptions in response to the threat of AI, *even if they both agree that AI is indeed a threat*. Future research should explore if this dynamic exists cross-nationally.

Methodologically, this study provides some of the first empirical tests of two new AI instruments, the AI Self Efficacy Scale (Wang, [2024](#)) and the EPOCH AI Risk and Augmentation dataset (Loaiza and Rigobon, [2024](#)), and finds individuals' political preferences interact with these metrics in logical, predictable ways when given information about the impact of AI on their incomes. In particular, our findings suggest that information might reprioritize or highlight personal income considerations among those in high AI-risk or high AI-augmentation occupations that were dormant in the absence of information.

Finally, these findings offer insight into the way the pre-distribution vs. redistribution dichotomy will translate into the AI-era. In Kuziemko et al.'s ([2023](#)) seminal paper, the authors show that the less educated in the US prefer "pre-distribution", that is regulating the source of economic disruption, while the more educated prefer traditional redistribution. I find that among pessimistic or economically vulnerable respondents, regulating AI is a more politically feasible demand than universal basic income (UBI) or tax-and-transfer expansion, but unlike Kuziemko et al. (2023), do *not* find that the more educated prefer redistribution in the context of AI. This distinction should be taken seriously by scholars and policymakers hoping to understand political reactions to technological change. This study suggests that voters may be far more receptive to

the politics of AI regulation than the redistributive, UBI focused politics of AI made famous by Andrew Yang in his 2020 bid for the American presidency.

Bibliography

- Acemoglu, Daron, and Simon Johnson. 2023. *Power and Progress: Our Thousand-Year Struggle over Technology and Prosperity*.
- Alder, Simeon, David Lagakos, and Lee E. Ohanian. 2014. "The Decline of the U.S. Rust Belt: A Macroeconomic Analysis." doi:[10.2139/ssrn.2586168](https://doi.org/10.2139/ssrn.2586168).
- Alesina, Alberto, and Paola Giuliano. "Preferences for Redistribution."
- "Anti-System Politics - Jonathan Hopkin - Oxford University Press."
<https://global.oup.com/academic/product/anti-system-politics-9780190699765?cc=us&lang=en&> (August 6, 2025).
- "Automation and the Welfare State: Technological Change as a Determinant of Redistribution Preferences - Stefan Thewissen, David Rueda, 2019."
<https://journals.sagepub.com/doi/abs/10.1177/0010414017740600> (August 11, 2025).
- Autor, David H., Frank Levy, and Richard J. Murnane. 2003. "The Skill Content of Recent Technological Change: An Empirical Exploration*." *The Quarterly Journal of Economics* 118(4): 1279–1333. doi:[10.1162/003355303322552801](https://doi.org/10.1162/003355303322552801).
- Benabou, Roland, and Efe A. Ok. 2001. "Social Mobility and the Demand for Redistribution: The Poup Hypothesis*." *The Quarterly Journal of Economics* 116(2): 447–87. doi:[10.1162/00335530151144078](https://doi.org/10.1162/00335530151144078).
- Berinsky, Adam J., and Gabriel S. Lenz. 2011. "Education and Political Participation: Exploring the Causal Link." *Political Behavior* 33(3): 357–73.
- Briel, Stephanie, Aderonke Osikominu, Gregor Pfeifer, Mirjam Reutter, and Sascha Satlukal. 2020. "Overconfidence and Gender Differences in Wage Expectations." *SSRN Electronic Journal*. doi:[10.2139/ssrn.3660264](https://doi.org/10.2139/ssrn.3660264).

- Broz, J. Lawrence, Jeffry Frieden, and Stephen Weymouth. 2021. "Populism in Place: The Economic Geography of the Globalization Backlash." *International Organization* 75(2): 464–94. doi:[10.1017/S0020818320000314](https://doi.org/10.1017/S0020818320000314).
- Bruni, Riccardo, Alessandro Gioffré, and Maria Marino. "In-Group Bias in Preferences for Redistribution: A Survey Experiment in Italy."
- Burke, Brian L., Spee Kosloff, and Mark J. Landau. 2013. "Death Goes to the Polls: A Meta-Analysis of Mortality Salience Effects on Political Attitudes." *Political Psychology* 34(2): 183–200. doi:[10.1111/pops.12005](https://doi.org/10.1111/pops.12005).
- Busemeyer, Marius R., and Alexander H. J. Sahm. 2022. "Social Investment, Redistribution or Basic Income? Exploring the Association Between Automation Risk and Welfare State Attitudes in Europe." *Journal of Social Policy* 51(4): 751–70. doi:[10.1017/S0047279421000519](https://doi.org/10.1017/S0047279421000519).
- Buser, Thomas, Gianluca Grimalda, Louis Putterman, and Joël van der Weele. 2020. "Overconfidence and Gender Gaps in Redistributive Preferences: Cross-Country Experimental Evidence." *Journal of Economic Behavior & Organization* 178: 267–86. doi:[10.1016/j.jebo.2020.07.005](https://doi.org/10.1016/j.jebo.2020.07.005).
- Cruces, Guillermo, Ricardo Pérez Truglia, and Martin Tetaz. "Biased Perceptions of Income Distribution and Preferences for Redistribution: Evidence from a Survey Experiment."
- Dinlersoz, Emin M., and Jeremy Greenwood. 2012. "The Rise and Fall of Unions in the U.S." doi:[10.3386/w18079](https://doi.org/10.3386/w18079).
- Ellis, Paul D. 2010. *The Essential Guide to Effect Sizes: Statistical Power, Meta-Analysis, and the Interpretation of Research Results*. Cambridge: Cambridge University Press. doi:[10.1017/CBO9780511761676](https://doi.org/10.1017/CBO9780511761676).
- Felten, Edward, Manav Raj, and Robert Seamans. 2021. "Occupational, Industry, and Geographic Exposure to Artificial Intelligence: A Novel Dataset and Its Potential Uses." *Strategic Management Journal* 42(12): 2195–2217. doi:[10.1002/smj.3286](https://doi.org/10.1002/smj.3286).

- “Field Experiments: Design, Analysis and Interpretation | Wiley.” *Wiley.com*.
<https://www.wiley.com/en-au/Field+Experiments%3A+Design%2C+Analysis+and+Interpretation-p-9780393979954> (August 7, 2025).
- Galipeau, Thomas. 2023. “The Impact of Political Memes: A Longitudinal Field Experiment.” *Journal of Information Technology & Politics* 20(4): 437–53.
 doi:[10.1080/19331681.2022.2150737](https://doi.org/10.1080/19331681.2022.2150737).
- Gelman, Andrew. 2007. “Rich State, Poor State, Red State, Blue State: What’s the Matter with Connecticut?” *Quarterly Journal of Political Science* 2(4): 345–67.
 doi:[10.1561/100.00006026](https://doi.org/10.1561/100.00006026).
- Golin, M., and C. Rauh. 2022. “The Impact of Fear of Automation.” *Janeway Institute Working Papers*. <https://ideas.repec.org/p/cam/camjip/2229.html> (August 6, 2025).
- Green, Jane, Zack Grant, Geoffrey Evans, and Gaetano Inglese. 2025. “Linking Artificial Intelligence Job Exposure to Expectations: Understanding AI Losers, Winners, and Their Political Preferences.” *Research & Politics* 12(2): 20531680251337897.
 doi:[10.1177/20531680251337897](https://doi.org/10.1177/20531680251337897).
- Hainmueller, Jens, Jonathan Mummolo, and Yiqing Xu. 2019. “How Much Should We Trust Estimates from Multiplicative Interaction Models? Simple Tools to Improve Empirical Practice.” *Political Analysis* 27(2): 163–92. doi:[10.1017/pan.2018.46](https://doi.org/10.1017/pan.2018.46).
- Heckert, Teresa M., Heather E. Droste, Patrick J. Adams, Christopher M. Griffin, Lisa L. Roberts, Michael A. Mueller, and Hope A. Wallis. 2002. “Gender Differences in Anticipated Salary: Role of Salary Estimates for Others, Job Characteristics, Career Paths, and Job Inputs.” *Sex Roles* 47(3): 139–51. doi:[10.1023/A:1021095005694](https://doi.org/10.1023/A:1021095005694).
- Heinrich, Tobias, and Christopher Witko. 2025a. “Self-Interest and Preferences for the Regulation of Artificial Intelligence.” *Journal of Information Technology & Politics* 22(3): 306–21. doi:[10.1080/19331681.2024.2370815](https://doi.org/10.1080/19331681.2024.2370815).
- Heinrich, Tobias, and Christopher Witko. 2025b. “Self-Interest and Preferences for the Regulation of Artificial Intelligence.” *Journal of Information Technology & Politics* 22(3): 306–21. doi:[10.1080/19331681.2024.2370815](https://doi.org/10.1080/19331681.2024.2370815).

- Imai, Kosuke, Luke Keele, and Dustin Tingley. 2010a. "A General Approach to Causal Mediation Analysis." *Psychological Methods* 15(4): 309–34. doi:[10.1037/a0020761](https://doi.org/10.1037/a0020761).
- Imai, Kosuke, Luke Keele, and Dustin Tingley. 2010b. "A General Approach to Causal Mediation Analysis." *Psychological Methods* 15(4): 309–34. doi:[10.1037/a0020761](https://doi.org/10.1037/a0020761).
- "Individual Risk Attitudes: Measurement, Determinants, and Behavioral Consequences | Journal of the European Economic Association | Oxford Academic." <https://academic.oup.com/jeea/article/9/3/522/2298422> (August 6, 2025).
- "Introducing ChatGPT." 2024. <https://openai.com/index/chatgpt/> (September 28, 2025).
- Iyengar, Shanto, and Masha Krupenkin. 2018. "Partisanship as Social Identity; Implications for the Study of Party Polarization." *The Forum* 16(1): 23–45. doi:[10.1515/for-2018-0003](https://doi.org/10.1515/for-2018-0003).
- "IZA Discussion Paper No. 912."
- Jacobs, Julian. 2024. "The Artificial Intelligence Shock and Socio-Political Polarization." *Technological Forecasting and Social Change* 199: 123006. doi:[10.1016/j.techfore.2023.123006](https://doi.org/10.1016/j.techfore.2023.123006).
- Kerrissey, Jasmine, and Evan Schofer. 2018. "Labor Unions and Political Participation in Comparative Perspective." *Social Forces* 97(1): 427–64. doi:[10.1093/sf/soy044](https://doi.org/10.1093/sf/soy044).
- Kim, Yeaji. 2023. "Absolutely Relative: How Education Shapes Voter Turnout in the United States." *Social Indicators Research*: 1–23. doi:[10.1007/s11205-023-03146-1](https://doi.org/10.1007/s11205-023-03146-1).
- Kuziemko, Ilyana, Nicolas Longuet Marx, and Suresh Naidu. "'Compensate the Losers?' Economic Policy and Partisan Realignment in the US."
- Kuziemko, Ilyana, Michael I. Norton, Emmanuel Saez, and Stefanie Stantcheva. 2015. "How Elastic Are Preferences for Redistribution? Evidence from Randomized Survey Experiments." *American Economic Review* 105(4): 1478–1508. doi:[10.1257/aer.20130360](https://doi.org/10.1257/aer.20130360).
- Lall, Ranjit. "The Political Economy of Artificial Intelligence: Evidence from Western Europe."

- Loaiza, Isabella, and Roberto Rigobon. “The EPOCH of AI: Human-Machine Complementarities at Work.”
- Loaiza, Isabella, and Roberto Rigobon. “The EPOCH of AI: Human-Machine Complementarities at Work.”
- Lown, Jean M. 2011. “2011 Outstanding AFCPE® Conference Paper: Development and Validation of a Financial Self-Efficacy Scale.” 22(2).
- MARES, ISABELA. 2000. “Strategic Alliances and Social Policy Reform: Unemployment Insurance in Comparative Perspective.” *Politics & Society* 28(2): 223–44. doi:[10.1177/0032329200028002004](https://doi.org/10.1177/0032329200028002004).
- McClure, Paul K. 2018. “‘You’re Fired,’ Says the Robot: The Rise of Automation in the Workplace, Technophobes, and Fears of Unemployment.” *Social Science Computer Review* 36(2): 139–56. doi:[10.1177/0894439317698637](https://doi.org/10.1177/0894439317698637).
- Mutz, Diana C. 2021. “Improving Experimental Treatments in Political Science.” In *Advances in Experimental Political Science*, eds. Donald P. Green and James N. Druckman. Cambridge: Cambridge University Press, 219–38. doi:[10.1017/9781108777919.016](https://doi.org/10.1017/9781108777919.016).
- “No Rage Against the Machines | Proceedings of the 2022 AAAI/ACM Conference on AI, Ethics, and Society.” *ACM Conferences*. doi:[10.1145/3514094.3534179](https://doi.org/10.1145/3514094.3534179).
- oecdstatistics. 2023. “Sizing up Welfare States: How Do OECD Countries Compare?” <https://oecdstatistics.blog/2023/02/02/sizing-up-welfare-states-how-do-oecd-countries-compare/> (August 7, 2025).
- Pierson, Paul. 2001. *The New Politics of the Welfare State*. OUP Oxford.
- Porter, Ethan, and Yamil R. Velez. 2022. “Placebo Selection in Survey Experiments: An Agnostic Approach.” *Political Analysis* 30(4): 481–94. doi:[10.1017/pan.2021.16](https://doi.org/10.1017/pan.2021.16).
- Rauh, Christopher, and Marta Golin. 2025. “The Impact of Fear of Automation.” doi:[10.2139/ssrn.5149312](https://doi.org/10.2139/ssrn.5149312).

- Ravallion, Martin, and Michael Lokshin. 2000. "Who Wants to Redistribute?: The Tunnel Effect in 1990s Russia." *Journal of Public Economics* 76(1): 87–104. doi:[10.1016/S0047-2727\(99\)00064-X](https://doi.org/10.1016/S0047-2727(99)00064-X).
- Schendstok, Matt, and Sydney Schreiner Wertz. "Occupational Exposure to Artificial Intelligence by Geography and Education."
- Spitz-Oener, Alexandra. 2006. "Technical Change, Job Tasks, and Rising Educational Demands: Looking Outside the Wage Structure." *Journal of Labor Economics* 24(2): 235–70. doi:[10.1086/499972](https://doi.org/10.1086/499972).
- Streeck, Wolfgang, and Anke Hassel. 2003. "Chapter 10: Trade Unions as Political Actors." <https://www.elgaronline.com/edcollchap/1840649798.00014.xml> (August 6, 2025).
- Sunshine Hillygus, D. 2005. "The MISSING LINK: Exploring the Relationship Between Higher Education and Political Engagement." *Political Behavior* 27(1): 25–47. doi:[10.1007/s11109-005-3075-8](https://doi.org/10.1007/s11109-005-3075-8).
- "The AI Workspace That Works for You." *Notion*. <https://www.notion.so> (August 13, 2025).
- "The Trade Origins of Economic Nationalism: Import Competition and Voting Behavior in Western Europe on JSTOR." <https://www.jstor.org/stable/26598793?seq=1> (August 6, 2025).
- Thewissen, Stefan, and David Rueda. 2019. "Automation and the Welfare State: Technological Change as a Determinant of Redistribution Preferences." *Comparative Political Studies* 52(2): 171–208. doi:[10.1177/0010414017740600](https://doi.org/10.1177/0010414017740600).
- Tomlinson, Kiran, Sonia Jaffe, Will Wang, Scott Counts, and Siddharth Suri. 2025. "Working with AI: Measuring the Occupational Implications of Generative AI." doi:[10.48550/arXiv.2507.07935](https://doi.org/10.48550/arXiv.2507.07935).
- Wallsten, Kevin J, and Gene Park. 2015a. "Confidence, Perception, and Politics in California: The Determinants of Attitudes toward Taxes by Level of Government." *California Journal of Politics and Policy* 7(2). doi:[10.5070/P2CJPP7225907](https://doi.org/10.5070/P2CJPP7225907).

- Wallsten, Kevin J, and Gene Park. 2015b. “Confidence, Perception, and Politics in California: The Determinants of Attitudes toward Taxes by Level of Government.” *California Journal of Politics and Policy* 7(2). doi:[10.5070/P2CJPP7225907](https://doi.org/10.5070/P2CJPP7225907).
- Wang, Yu-Yin, and Yu-Wei Chuang. 2024. “Artificial Intelligence Self-Efficacy: Scale Development and Validation.” *Education and Information Technologies* 29(4): 4785–4808. doi:[10.1007/s10639-023-12015-w](https://doi.org/10.1007/s10639-023-12015-w).
- Webb, Michael. “The Impact of Artificial Intelligence on the Labor Market.”
- Zhang, Baobao. 2022. “No Rage Against the Machines: Threat of Automation Does Not Change Policy Preferences.” doi:[10.2139/ssrn.3455501](https://doi.org/10.2139/ssrn.3455501).

Appendix

Survey Text

Survey Questions

White-Collar Status

- Do you consider your current job to be “white-collar”?
 - Yes
 - No
 - Not sure

Occupation

- What is your current or most recent occupation? (*Dropdown autocomplete of SOC 2018 Coded Occupations*)

Partisanship

- Generally speaking, do you usually think of yourself as:
 - Republican
 - Democrat
 - Independent
 - Other
 - No preference

Education

- What is the highest level of education you have completed?
 - Some high school or less
 - High school diploma or GED
 - Some college, but no degree
 - Associate's or technical degree
 - Bachelor's degree
 - Graduate or professional degree
 - Prefer not to say

Income

- What is your average monthly income? (*Free response*)

Childhood Socioeconomic Status

- How would you describe your socioeconomic class growing up?
 - Upper class
 - Upper middle class
 - Middle class
 - Lower middle class
 - Lower class
 - I don't know

Future Socioeconomic Expectations

- What socioeconomic class do you expect to be in the future?
 - Upper class
 - Upper middle class
 - Middle class
 - Lower middle class
 - Lower class
 - I don't know

AI Self-Efficacy (AISES)

How true are the following statements?

1. When interacting with AI technologies/products, I feel very relaxed.
2. I can happily interact with AI technologies/products smoothly.

3. AI technologies/products jargon does not baffle me.
4. AI technologies/products are good aids to learning.
5. AI technologies/products help me to save a lot of time.
6. I find it easy to get AI technologies/products to do what I want it to do.

Response options:

- ☐ Exactly true
- ☐ Moderately true
- ☐ Hardly true
- ☐ Not at all true
- ☐ Not sure

Risk Aversion

- On a scale from 0 to 10, how willing are you to take risks?
(0 = not at all willing, 10 = very willing)

Attention Check

- This is an attention check. Please select **“Moderately important”**.
 - ☐ Not at all important
 - ☐ Slightly important
 - ☐ Moderately important
 - ☐ Very important
 - ☐ Extremely important

AI Knowledge Questions

1. What is OpenAI?
 - ☐ A tech company
 - ☐ An organization lobbying for open borders
 - ☐ An organization lobbying for nondiscriminatory elections
 - ☐ An organization lobbying for open access to government documents
 - ☐ Other / I don't know
2. Large language models like ChatGPT and Claude are primarily trained to do what?
 - ☐ Predict the next word in a sequence of text
 - ☐ Translate only between languages
 - ☐ Search the internet in real time
 - ☐ Perform mathematical proofs
3. What does “RLHF” stand for?
 - ☐ Reinforcement Learning from Human Feedback
 - ☐ Risky Lifestyle Health Factors
 - ☐ Renewable Liquid Hydrogen Fuel

- Other / I don't know
4. Who is Sam Altman?
- Influential technology entrepreneur
 - Trump cabinet member
 - NCAA football star
 - Screenwriter
 - Other / I don't know

Treatment Conditions

Participants were randomly assigned to either a placebo control video or an AI framing treatment (Complementarity or Substitution). All conditions were followed by an attention check:

- **Complementarity Video:** Described how AI helps white-collar workers be more productive and earn more.
- **Substitution Video:** Described how AI is replacing white-collar jobs.

Video Comprehension Check Example:

“What was the video you just watched about?”

- How AI helps white collar workers
- The Grand Canyon
- The Pyramids of Giza
- How AI helped discover bacteria (etc., depending on condition)

Mediators

AI Unemployment Concern

- What is the chance you estimate you will be unemployed due to AI in the next 5 years?
 - Very likely
 - Somewhat likely
 - No effect
 - Somewhat unlikely
 - Very unlikely

Expected Personal Income Change

- Estimate the percentage impact AI will have on your income. (-100 to $+100$)

Sociotropic AI Impact

- What effect will AI have on the total number of jobs in the economy?
 - Greatly decrease
 - Somewhat decrease
 - No major effect
 - Somewhat increase
 - Greatly increase

Redistribution Preferences

General Redistribution

Some people feel that the government should increase taxes a lot and spend much more on health and social services such as unemployment benefits and low-cost healthcare programs like Medicaid. Other people feel that the government should cut taxes a lot and spend much less on health and social services. And other people have views somewhere in-between.

- Please use the slider to indicate where your views fall on this scale.
(0 = Cut taxes a lot and spend less on social services; 100 = Increase taxes a lot and spend more)

Tax the Rich

- What should happen to tax rates on income over \$400k? (Slider: 0 = lower, 100 = raise)

Tax Corporations

- What should happen to tax rates on large tech companies? (Slider: 0 = lower, 100 = raise)

Job Retraining Support

- The government should fund retraining for workers who lose jobs to AI.
◦ Strongly agree → Strongly disagree (5-point scale)

Union Interest

- How interested are you in joining a labor union?
◦ Strongly interested → Not at all interested
◦ I am already a member of a union

Universal Basic Income

- If UBI were implemented, how much should each person receive monthly? (0–10,000 dollars)

Regulation Preferences

Ban Layoffs

- The government should stop companies from replacing workers with AI.
◦ Strongly agree → Strongly disagree

Mandate Safety Measures

- The government should require safety measures for new technologies.
◦ Strongly agree → Strongly disagree

New AI Agency

- Support for a new federal agency to license AI projects.
 - Strongly agree → Strongly disagree

Overall Regulation Level

- Should the federal government do more or less to regulate AI development?
 - A lot more → A lot less
 - Not sure

Manipulation Check Details by Party

Democrats			
	Expected % Income Impact from AI	Expected Impact of AI on Jobs (standardized)	AI Unemployment Concern (standardized)
Substitution	-1.037 (3.611)	-0.322** (0.100)	-0.079 (0.124)
Complementarity	8.948* (3.562)	0.143 (0.119)	-0.077 (0.119)
Num.Obs.	431	430	431
RMSE	29.62	0.93	1.03
+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001			
Mediator Effects on Democrats			

Republicans			
	Expected % Income Impact from AI	Expected Impact of AI on Jobs (standardized)	AI Unemployment Concern (standardized)
Substitution	-11.405* (4.611)	-0.446*** (0.124)	-0.100 (0.130)
Complementarity	-1.863 (4.615)	-0.006 (0.138)	-0.114 (0.126)
Num.Obs.	373	373	373
RMSE	36.01	1.02	1.01
+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001			
Mediator Effects on Republicans			


Risk Aversion and Treatment Effects

Respondents were asked how comfortable they are taking risks on a 0-10 sliding scale, where a value of 0 indicates not willing to take risks at all and a value of 10 indicates being very willing to take risks. Higher values therefore reflect greater willingness to take risks.

Risk aversion was pre-registered as a hypothesized moderator: those less willing to take risks were expected to show stronger support for regulation and redistribution, consistent with a preference for labor protections and safety nets. Conversely, risk-tolerant individuals were expected to show less support for such policies.

I largely find evidence that I can reject the null hypothesis. In keeping with expectations, I first test the relationship between risk aversion and our outcomes of interest among those in the control group to understand how this personal quality impacts policy preferences without treatment cues. Among our survey respondents, I find that being *more* willing to take risks is associated with less support for general redistribution, less support for AI regulation, less support for taxing the rich, and less support for taxing corporations. Inconsistent with expectations is that being more willing to take risks is associated with more support for UBI, though it is possible this has something to do with the somewhat fringe nature of this policy position.

How does job-threat framing interact with the willingness to take risks? No significant interactions emerge in the complementarity treatment group. However, the substitution treatment appears to mitigate the effect of risk attitudes on tax preferences ($p < .05$). Specifically, the negative association between risk tolerance and support for taxing the rich and corporations is substantially weakened. This suggests that substitution framing, which highlights job replacement risks from AI, may heighten personal economic concerns, making even risk-tolerant individuals more receptive to taxing the wealthy.



	redistribution_pref_z	ubi_support_z	job_retraining_z	regulation_index_z	regulation_agency_z	tax_the_rich_z	tax_corporations_z	union_join_intent_z
Substitution	0.094	-0.142+	-0.102	0.116	0.120	0.138+	0.124	-0.000
	(0.077)	(0.077)	(0.077)	(0.075)	(0.075)	(0.077)	(0.076)	(0.078)
Complementarity	0.028	-0.070	-0.081	0.025	-0.029	0.062	-0.025	-0.028
	(0.077)	(0.078)	(0.074)	(0.077)	(0.079)	(0.074)	(0.076)	(0.079)
Risk Aversion (z)	-0.143*	0.126*	-0.012	-0.103+	-0.017	-0.237***	-0.238***	0.099+
	(0.057)	(0.059)	(0.051)	(0.054)	(0.052)	(0.055)	(0.051)	(0.054)
Subst. × Risk	0.127	-0.062	-0.002	-0.033	-0.109	0.163*	0.199**	0.023
	(0.080)	(0.081)	(0.081)	(0.077)	(0.072)	(0.080)	(0.075)	(0.079)
Compl. × Risk	0.040	0.086	-0.069	-0.085	-0.074	-0.035	0.048	-0.103
	(0.085)	(0.085)	(0.076)	(0.079)	(0.081)	(0.081)	(0.077)	(0.082)
Num.Obs.	993	981	991	991	990	992	992	957
RMSE	0.99	0.99	1.00	0.99	0.99	0.98	0.98	0.99


+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001

AI Self-Efficacy Robustness Checks

Here, treatment effects are displayed when AI Self-Efficacy is dichotomized such that a value of 2 or less (on the 1-4 Likert scale) is considered low self efficacy, rather than dichotomized as above or below the median.

Treatment Effects by AI Self-Efficacy

● $p < .05$, ○ n.s.




How Does AI Risk Interact Vary by Subgroups of Interest?

Among respondents with low AI self-efficacy, I observe distinct moderating patterns. The substitution treatment showed a significant positive interaction with AI risk for union interest ($\beta = 2.03$, $p = 0.042$), suggesting that low-efficacy workers in high-risk occupations became more interested in unionization when exposed to substitution framing. The complementarity treatment exhibited a significant positive interaction with AI risk for support of taxing the rich ($\beta = 2.02$, p


= 0.019). When examining AI augmentation as the moderator, the complementarity treatment demonstrated significant negative interactions for both UBI support ($\beta = -4.92$, $p = 0.009$) and taxing the rich ($\beta = -4.79$, $p = 0.005$), indicating that low-efficacy workers in jobs with high augmentation potential became less supportive of these redistributive policies when exposed to complementarity framing.

Significant Only - LowSelfEfficacy | Complementarity × Augmentation




The pattern reversed among high AI self-efficacy respondents. Here, the substitution treatment produced a significant negative interaction with AI risk for support of taxing the rich ($\beta = -1.42$, $p = 0.036$), while the complementarity treatment showed a significant positive interaction with AI risk for support of taxing corporations ($\beta = 2.44$, $p < 0.001$). For AI augmentation, only one significant effect emerged: the complementarity treatment exhibited a negative interaction with augmentation for support of taxing corporations ($\beta = -3.32$, $p = 0.014$). These divergent patterns suggest that AI self-efficacy moderates how workers in different occupational contexts respond to framing about AI's labor market effects.

Significant Only - HighSelfEfficacy | Complementarity × Augmentation



When subset by partisanship, no significant interactions are found between AI risk and policy preferences, with the exception of Republicans and support for taxing the rich. Republicans in the control group are more supportive of taxing the rich as AI risk goes up. Interestingly, that trend is *reversed* when primed with the substitution treatment.

Significant Only - Republicans | Substitution × Risk



In all, I largely find support for AI risk and AI augmentation moving tax policy in expected directions, but do not find statistically significant results for other outcomes of interest (regulation, job retraining, redistribution, etc.).

Placebo videos were checked to ensure that the 4 different placebo videos respondents were randomly assigned adequately captured the nonspecific effects (NSEs) I aimed to associate with the control group (Porter and Velez 2021). This was done using a one-way ANOVA analysis, and a full table showing these results are in Appendix table A.1.

A.1 Placebo Comparison Test

A.2: Demographic Approximations of White-Collar Workers age 25-54

In order to run this experiment on a representative sample of individuals who are the core audience for messaging regarding AI's potential to either complement or substitute their labor power, I approximate the demographic breakdown of white-collar workers from ages 25-54 so I can then set quotas for that demographic breakdown with Prolific. The US Census provides the Equal Employment Opportunity tabulation (most recently for 2013-2018) on occupational groups by Sex and Ethnicity based on subfields of employment¹⁴. I consider the italicized categories white-collar for the purposes of this analysis. I use quotas based on demographics of people 16 and up, since racial and gender breakdown of white collar workers ages 25 to 54 is not readily available.¹⁵

White Collar

- *Science, engineering and computer professionals*
- *Healthcare practitioner professionals*
- *Other professional workers*
- *Technicians*
- *Sales workers*
- *Administrative support workers*

Non White-Collar

- Construction and extractive craft workers
- Installation, maintenance and repair craft workers
- Production operative workers
- Transportation and material moving operative workers
- Laborers and helpers
- Protective service workers

¹⁴<https://www.census.gov/acs/www/data/eoo-data/eoo-tables-2018/tableview.php?geotype=nation&usVal=us&filetype=all3w&geoName=United%20States>

¹⁵ "I follow the EEO scheme: Hispanic or Latino (any race); and among non-Hispanics: White alone, Black alone, AIAN alone, Asian alone, NH/PI alone, and a residual 'Non-Hispanic Multiracial/Other' ('balance') comprising non-Hispanic respondents of two or more races and 'some other race'.

- Service workers, except protective

White-Collar Demographics (EEO-ALL03W, 2014–2018, Age 16+)

Dimension	Group	Count	Percent
Race/Ethnicity	Hispanic or Latino (any race)	9,832,660	11.7%
Race/Ethnicity	White alone (Not Hispanic)	58,502,815	69.5%
Race/Ethnicity	Black or African American alone (Not Hispanic)	7,878,330	9.4%
Race/Ethnicity	Asian alone (Not Hispanic)	5,757,095	6.8%
Race/Ethnicity	American Indian/Alaska Native alone (Not Hispanic)	356,000	0.4%
Race/Ethnicity	Native Hawaiian/Other Pacific Islander alone (Not Hispanic)	113,030	0.1%
Race/Ethnicity	Balance of Not Hispanic (e.g., Two+ races)	1,683,635	2.0%
Sex	Male	37,529,320	44.6%
Sex	Female	46,594,245	55.4%

Prolific Survey Demographics

Table 1: Gender Distribution

CATEGORY	COUNT	%
Female	557	54.70%
Male	462	45.30%


Table 2: Political Affiliation

CATEGORY	COUNT	%
----------	-------	---

Democrat	423	41.50%
Independent	186	18.30%
Other	43	4.20%
Republican	367	36.00%

Table 3: Employment Status

CATEGORY	COUNT	%
Full-Time	852	83.60%
Part-Time	167	16.40%



Descriptive Statistics:

Ethnicity Distribution of Survey Respondents

CATEGORY	COUNT	%
Black/African American	93	9.10%
East Asian	22	2.20%
Latino/Hispanic	121	11.90%
Mixed	10	1.00%
Native American or Alaskan Native	7	0.70%

Other (please feel free to let us know your ethnicity via email)	7	0.70%
South Asian	17	1.70%
South East Asian	18	1.80%
White	693	68.00%
NA	31	3.00%

Table: AI Knowledge Index (0–4 correct)

Metric	Value
N	995.00
Mean	2.74
SD	1.08
Median	3.00
Min	0.00
Max	4.00